

MA615 Assignment 2

Andrew Zhang

9/23/2018

R for Data Science Exercises

3.5.1

Question 2.

What do the empty cells in plot with *facet_gird(drv ~ cyl)* mean? How do they relate to this plot?

The empty cells mean there are no rows with that specific combination within the dataset. In relation to this plot, if there are empty cells, that just means that mpg doesn't contain the specific combinations of variables being asked for.

Question 3.

What plots does the following code make? What does . do?

The first plot generates highway miles per gallon vs. engine displacement, separating the three levels within the various types of drive options(front wheel drive(f), rear wheel drive(r), and 4 wheel drive(4)). The second plot generates engine displacement vs. highway miles per gallon, separating the four levels within the various types of number of cylinders(4, 5, 6, 8). The . operator is a placeholder, allowing for only one dimension when dealing with multiple variables.

3.6.1

Question 6.

Recreate the R code necessary to generate the following graphs.


```
install.packages("ggplot2", repos = "https://cran.r-project.org")

##
## The downloaded binary packages are in
## /var/folders/9r/b1y49xdd2mg853hg98c6c6780000gn/T//RtmpXHcg3v/downloaded_packages
library(ggplot2)

## Warning: package 'ggplot2' was built under R version 3.4.4


data(mpg)
ggplot(mpg) +
  geom_point(aes(x = displ, y = hwy)) +
  geom_smooth(aes(x = displ, y = hwy), se = FALSE)

## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```




```
ggplot(mpg, aes(x = displ, y = hwy, group = drv)) +  
  geom_point() + geom_smooth(se = FALSE)
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```




```
ggplot(mpg, aes(x = displ, y = hwy, group = drv)) +  
  geom_point(aes(colour = drv)) + geom_smooth(se = FALSE)
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```




```
ggplot(mpg, aes(x = displ, y = hwy, group = drv)) +  
  geom_point(aes(colour = drv)) + geom_smooth(aes(colour = drv), se = FALSE)
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```




```
ggplot(mpg, aes(x = displ, y = hwy, group = drv)) +  
  geom_point(aes(colour = drv)) + geom_smooth(aes(linetype = drv), se = FALSE)
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```



```
ggplot(mpg, aes(x = displ, y = hwy, group = drv)) +  
  geom_point(size = 4, colour = "white") + geom_point(aes(colour = drv), se = FALSE)
```

Warning: Ignoring unknown parameters: se



5.2.4

Question 1.

```
install.packages("nycflights13", repos = "https://cran.r-project.org")

##
## The downloaded binary packages are in
## /var/folders/9r/b1y49xdd2mg853hg98c6c6780000gn/T//RtmpXHcg3v/downloaded_packages
library(nycflights13)

## Warning: package 'nycflights13' was built under R version 3.4.4
library(dplyr)

## Warning: package 'dplyr' was built under R version 3.4.4
##
## Attaching package: 'dplyr'
## The following objects are masked from 'package:stats':
##
##   filter, lag
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
attach(flights)
```

```
## Had an arrival delay of two or more hours
```

```
filter(flights, arr_delay >= 120)
```

```
## Warning: package 'bindrcpp' was built under R version 3.4.4
```

```
## # A tibble: 10,200 x 19
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     1     811           630          101    1047
## 2  2013     1     1     848          1835          853    1001
## 3  2013     1     1     957           733          144    1056
## 4  2013     1     1    1114           900          134    1447
## 5  2013     1     1    1505          1310          115    1638
## 6  2013     1     1    1525          1340          105    1831
## 7  2013     1     1    1549          1445           64    1912
## 8  2013     1     1    1558          1359          119    1718
## 9  2013     1     1    1732          1630           62    2028
##10  2013     1     1    1803          1620          103    2008
```

```
## # ... with 10,190 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

```
## Flew to Houston
```

```
filter(flights, dest == "IAH" | dest == "HOU")
```

```
## # A tibble: 9,313 x 19
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     1     517           515           2     830
## 2  2013     1     1     533           529           4     850
## 3  2013     1     1     623           627          -4     933
## 4  2013     1     1     728           732          -4    1041
## 5  2013     1     1     739           739           0    1104
## 6  2013     1     1     908           908           0    1228
## 7  2013     1     1    1028          1026           2    1350
## 8  2013     1     1    1044          1045          -1    1352
## 9  2013     1     1    1114           900          134    1447
##10  2013     1     1    1205          1200           5    1503
```

```
## # ... with 9,303 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

```
## Were operated by United, American, or Delta
```

```
filter(flights, carrier == "UA" | carrier == "AA" | carrier == "DL")
```

```
## # A tibble: 139,504 x 19
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     1     517           515           2     830
## 2  2013     1     1     533           529           4     850
## 3  2013     1     1     542           540           2     923
## 4  2013     1     1     554           600          -6     812
```



```
## 5 2013 1 1 554 558 -4 740
## 6 2013 1 1 558 600 -2 753
## 7 2013 1 1 558 600 -2 924
## 8 2013 1 1 558 600 -2 923
## 9 2013 1 1 559 600 -1 941
## 10 2013 1 1 559 600 -1 854
## # ... with 139,494 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

```
## Departed in summer(July, August, September)
filter(flights, month >= 7 & month <= 9)
```

```
## # A tibble: 86,326 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1 2013     7     1       1         2029          212     236
## 2 2013     7     1       2         2359           3     344
## 3 2013     7     1      29         2245         104     151
## 4 2013     7     1     43         2130         193     322
## 5 2013     7     1     44         2150         174     300
## 6 2013     7     1     46         2051         235     304
## 7 2013     7     1     48         2001         287     308
## 8 2013     7     1     58         2155         183     335
## 9 2013     7     1    100         2146         194     327
## 10 2013     7     1    100         2245         135     337
## # ... with 86,316 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

```
## Arrived more than two hours late, but didn't leave late
filter(flights, arr_delay > 120 & dep_delay <= 0)
```

```
## # A tibble: 29 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1 2013     1    27    1419         1420          -1    1754
## 2 2013    10     7    1350         1350           0    1736
## 3 2013    10     7    1357         1359          -2    1858
## 4 2013    10    16     657          700          -3    1258
## 5 2013    11     1     658          700          -2    1329
## 6 2013     3    18    1844         1847          -3     39
## 7 2013     4    17    1635         1640          -5    2049
## 8 2013     4    18     558          600          -2    1149
## 9 2013     4    18     655          700          -5    1213
## 10 2013     5    22    1827         1830          -3    2217
## # ... with 19 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

```
## Were delayed by at least an hour, but made up over 30 minutes in flight
filter(flights, dep_delay >= 60 & dep_delay - arr_delay > 30)
```

```
## # A tibble: 1,844 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
## 1  2013     1     1    2205           1720        285     46
## 2  2013     1     1    2326           2130        116    131
## 3  2013     1     3    1503           1221        162   1803
## 4  2013     1     3    1839           1700         99   2056
## 5  2013     1     3    1850           1745         65   2148
## 6  2013     1     3    1941           1759        102   2246
## 7  2013     1     3    1950           1845         65   2228
## 8  2013     1     3    2015           1915         60   2135
## 9  2013     1     3    2257           2000        177     45
## 10 2013     1     4    1917           1700        137   2135
## # ... with 1,834 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

```
## Departed between midnight and 6AM(inclusive)
filter(flights, dep_time <= 600 | dep_time == 2400)
```

```
## # A tibble: 9,373 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
## 1  2013     1     1     517           515         2     830
## 2  2013     1     1     533           529         4     850
## 3  2013     1     1     542           540         2     923
## 4  2013     1     1     544           545        -1    1004
## 5  2013     1     1     554           600        -6     812
## 6  2013     1     1     554           558        -4     740
## 7  2013     1     1     555           600        -5     913
## 8  2013     1     1     557           600        -3     709
## 9  2013     1     1     557           600        -3     838
## 10 2013     1     1     558           600        -2     753
## # ... with 9,363 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

Question 2.

Another useful dplyr filtering helper is *between()*. What does it do? Can you use it simplify the code needed to answer the previous challenges?

between() is a simplified version of `x >= left & x <= right`.

```
filter(flights, between(month, 7, 9))
```

```
## # A tibble: 86,326 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
## 1  2013     7     1     1           2029        212    236
## 2  2013     7     1     2           2359         3    344
## 3  2013     7     1    29           2245        104    151
## 4  2013     7     1    43           2130        193    322
```

```
## 5 2013 7 1 44 2150 174 300
## 6 2013 7 1 46 2051 235 304
## 7 2013 7 1 48 2001 287 308
## 8 2013 7 1 58 2155 183 335
## 9 2013 7 1 100 2146 194 327
## 10 2013 7 1 100 2245 135 337
## # ... with 86,316 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>

filter(flights, !between(dep_time, 601, 2359))
```

```
## # A tibble: 9,373 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1 2013     1     1     517           515           2     830
## 2 2013     1     1     533           529           4     850
## 3 2013     1     1     542           540           2     923
## 4 2013     1     1     544           545          -1    1004
## 5 2013     1     1     554           600          -6     812
## 6 2013     1     1     554           558          -4     740
## 7 2013     1     1     555           600          -5     913
## 8 2013     1     1     557           600          -3     709
## 9 2013     1     1     557           600          -3     838
## 10 2013     1     1     558           600          -2     753
## # ... with 9,363 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

Question 3.

How many flights have a missing dep_time? What other variables are missing? What might these rows represent?

```
summary(flights)

## Warning in as.POSIXlt.POSIXct(x, tz): unknown timezone 'zone/tz/2018e.1.0/'
## zoneinfo/America/New_York'

##      year      month      day      dep_time
## Min.   :2013   Min.   : 1.000   Min.   : 1.00   Min.   : 1
## 1st Qu.:2013   1st Qu.: 4.000   1st Qu.: 8.00   1st Qu.: 907
## Median :2013   Median : 7.000   Median :16.00   Median :1401
## Mean   :2013   Mean   : 6.549   Mean   :15.71   Mean   :1349
## 3rd Qu.:2013   3rd Qu.:10.000   3rd Qu.:23.00   3rd Qu.:1744
## Max.   :2013   Max.   :12.000   Max.   :31.00   Max.   :2400
##                                     NA's   :8255
## sched_dep_time  dep_delay      arr_time  sched_arr_time
## Min.   : 106   Min.   : -43.00   Min.   : 1     Min.   : 1
## 1st Qu.: 906   1st Qu.: -5.00    1st Qu.:1104   1st Qu.:1124
## Median :1359   Median : -2.00    Median :1535   Median :1556
## Mean   :1344   Mean   : 12.64    Mean   :1502   Mean   :1536
## 3rd Qu.:1729   3rd Qu.: 11.00    3rd Qu.:1940   3rd Qu.:1945
```

```
## Max. :2359 Max. :1301.00 Max. :2400 Max. :2359
## NA's :8255 NA's :8713
## arr_delay carrier flight tailnum
## Min. : -86.000 Length:336776 Min. : 1 Length:336776
## 1st Qu.: -17.000 Class :character 1st Qu.: 553 Class :character
## Median : -5.000 Mode :character Median :1496 Mode :character
## Mean : 6.895 Mean :1972
## 3rd Qu.: 14.000 3rd Qu.:3465
## Max. :1272.000 Max. :8500
## NA's :9430
## origin dest air_time distance
## Length:336776 Length:336776 Min. : 20.0 Min. : 17
## Class :character Class :character 1st Qu.: 82.0 1st Qu.: 502
## Mode :character Mode :character Median :129.0 Median : 872
## Mean :150.7 Mean :1040
## 3rd Qu.:192.0 3rd Qu.:1389
## Max. :695.0 Max. :4983
## NA's :9430
## hour minute time_hour
## Min. : 1.00 Min. : 0.00 Min. :2013-01-01 05:00:00
## 1st Qu.: 9.00 1st Qu.: 8.00 1st Qu.:2013-04-04 13:00:00
## Median :13.00 Median :29.00 Median :2013-07-03 10:00:00
## Mean :13.18 Mean :26.23 Mean :2013-07-03 05:22:54
## 3rd Qu.:17.00 3rd Qu.:44.00 3rd Qu.:2013-10-01 07:00:00
## Max. :23.00 Max. :59.00 Max. :2013-12-31 23:00:00
##
```

There are 8255 missing flights for *dep_time*, 8255 missing flights for *dep_delay*, 8713 missing flights for *arr_time*, 9430 missing flights for *arr_delay*, and 9430 missing flights for *air_time*. It is possible that these flights weren't able to depart due to circumstances. It is also possible that the data was simply lost.

Question 4.

Why is $NA \wedge 0$ not missing? Why is $NA / TRUE$ not missing? Why is $FALSE \& NA$ not missing? Can you figure out the general rule?

$NA \wedge 0$ equals 1 since anything raised to 0 is 1. Therefore, $NA \wedge 0$ cannot be missing if it equals 1. $NA / TRUE$ will return *TRUE* if either side of the $/$ is true. Therefore, $NA / TRUE$ will never evaluate to missing. $FALSE \& NA$ will return *TRUE* or *FALSE* if the operation is proven to be true or false. Therefore, it will never return as missing. $NA \emptyset$ is interesting in that NA can represent *Inf*. If we were to multiply *Inf* and 0, we would get *NaN* not *NA*.